Gonzalo Díaz Amor

Data Scientist | GenAI & MLOps Specialist | LLM & RAG Engineer

gdiazamor.ds@gmail.com | +34693299747 | LinkedIn | GitHub | Portfolio

---

Professional Summary

Data Scientist with 5+ years of end-to-end experience across the full ML lifecycle — from statistical modeling and causal inference to LLM-based architectures, RAG pipelines, and MLOps. Background spans pharma, healthtech, and mobility, with production deployments on AWS, GCP, and Azure. Currently designing and scaling an enterprise LLM-as-a-Judge evaluation system, including architecture design, bias-mitigation strategies, and automated BOCR audit pipelines. Proven ability to deliver in fast-moving environments: built data lakes integrating multi-hospital sources, shipped FastAPI-based ML inference services, and led clinical programming for EMA-submission studies. Equally comfortable owning a full pipeline solo or collaborating cross-functionally with research and engineering teams.

Relevant Experience

AI/ML Research Engineer

BeWay | *Jan 2026 - Present*

- Led design and deployment of the LLM-as-a-Judge system for the RNA audit pipeline, defining evaluation taxonomy and multi-scenario scoring workflows.
- Scaled automated fidelity, pairwise, and holistic ranking checks across BOCR artifacts to improve review consistency and reduce manual validation.
- Built reusable Python/R tooling and AWS orchestration to run audit evaluations at scale for multiple model and prompt variants.
- *Tech stack:* Python, R, AWS (S3, Bedrock, SageMaker, Lambda, ECR), Docker, Git, CI/CD.

Behavioral Data Scientist

BeWay | *Jun 2025 - Oct 2025*

- Led Geo Data Science analysis across hundreds of features to measure product penetration over 60M+ tickets and inform strategic decisions.
- Owned survival and funnel analytics on 10k+ customers to identify conversion barriers and accelerate time-to-first-purchase.
- Transitioned into the current AI research role after four months of cross-functional delivery to scale advanced LLM and RAG systems.
- *Tech stack:* R, Python, Quarto, AWS (S3, Athena, Redshift), PostgreSQL, Git.

AI/ML Specialist Data Scientist

Daiichi-Sankyo | *Oct 2025 - Dec 2025*

- Delivered a 3-month Causal AI engagement identifying environmental drivers across hundreds of variables for patient outcomes using Double ML and Causal Forest.
- Built the data infrastructure and high-performance analytics scripts for Google Cloud deployment.

- Originally a 12-month engagement shortened due to internal budget restructuring; departed on excellent terms with strong professional references.
- *Tech stack:* Python, Google Cloud.

Data Science Specialist

Minsait | *Nov 2023 - Jun 2025*

- Built an enterprise OMOP-CDM data lake integrating clinical and operational hospital sources for analytics readiness.
- Owned RESTful ML/AI inference APIs with FastAPI, enabling embeddings, vector search, and LLM-powered analytics for downstream applications.
- Deployed NLP/ML models in production with Docker and Airflow, standardizing end-to-end workflows and reducing manual handoffs.
- *Tech stack:* R, Python, Docker, FastAPI, PowerShell, GitLab, MySQL, PostgreSQL, Spark, Airflow, OHDSI.

Data Scientist

TFS HealthScience | *Jun 2022 - May 2023*

- Served as programming lead for a Phase IV study, validating SDTM and ADaM workflows and saving the team 3+ years of research effort.
- Built regularized logistic regression models to improve predictive performance and support regulatory submission.
- Delivered propensity-matched comparative analyses and regulatory-ready statistical plans and study reports.
- *Tech stack:* R, SAS, Python, CDISC, OHDSI.

Data Scientist

Animal Data Analytics | *Sep 2020 - Jun 2022*

- Delivered analytics and ML solutions for 5+ international clients across clinical and commercial engagements.
- Led a small team to build a regression model that improved predictive performance by 7% and translated results into RMarkdown reports.
- Applied survival analysis methods to time-to-event outcomes and presented actionable insights to stakeholders.
- *Tech stack:* R, Python, SAS, RMarkdown.

Education

- **MSc. Big Data & Data Science** | VIU – Valencia, Remote (2025)
  - Advanced training in large-scale data processing, cloud architectures, and AI/ML systems.
- **BSc. Statistics** | UVa – Valladolid, Spain (2020)
  - Strong foundation in statistical modeling, probability theory, and quantitative methods.

Notable Additional Education (see LinkedIn for more):

- **Explainable AI (XAI)** | DPHi (2023)

- **Biomedical Image Analysis in Python** | DataCamp (2023)
- **Deep Learning Online Bootcamp** | DPHi (2021)
- **Visual Studio Code for Pharma R** | R in Pharma (2023)
- **Data Analysis with Python** | freeCodeCamp (2023)
- **Continuous Learning in Data Science** | Various platforms

---

Certifications

- **AI Agents Intensive Course** — Google x Kaggle *(2025)*
  - Hands-on training designing and deploying AI agent systems with LLMs, tools, memory, and retrieval.
- **6.86x: Machine Learning with Python – From Linear Models to Deep Learning** — MITx *(2021)*
  - Deep dive into linear models, deep learning, and reinforcement learning with hands-on Python projects.
- **LangChain for LLM Development** — freeCodeCamp *(2023)*
  - Practical LLM development with embeddings, prompts, vector search, and generative workflow orchestration.
- **Google Data Analytics Certificate** — Google *(2023)*
  - Foundational analytics training in R, SQL, Python, and visualization.

Technical Skills

- **Programming Languages:** Python (LLM APIs, LangChain, FastAPI), R, SQL, PowerShell.
- **GenAI & LLMs:** RAG pipelines, embeddings, vector search, prompt engineering, LangChain, LLM-based architectures.
- **ML & Data Science:** PyTorch, TensorFlow, Keras, Scikit-Learn, Pandas, NumPy, Causal ML (Causal Forest, Double ML), survival analysis, predictive modeling.
- **MLOps & Cloud:** Docker, Apache Airflow, AWS, Azure, GCP, Databricks, PySpark, Hadoop.
- **Data Engineering:** ETL pipelines, FastAPI, OMOP-CDM (OHDSI), WhiteRabbit, Usagi, CDISC standards.
- **Databases & BI:** MySQL, PostgreSQL, MongoDB, Cassandra, PowerBI, Tableau.
- **Version Control & Collaboration:** Git, GitHub, GitLab.
- **Languages:** Spanish (Native), English (Advanced).

References and recommendation letters available upon request.